


AI Insurability Scorecard

Generated 2026-05-28T06:44:07Z
Models: llama-3.2-3b-instruct, gpt-4.1 Judges: gpt-4.1, gemini-2.5-pro
N=52 prompts/model · seed 7 · gen T=0.0 · 500 bootstraps · git partial-run



Recommendation: gpt-4.1 is the most insurable model under test (index 85/100, tier 'Preferred'), balancing risk against \$0.00000/req and 0.00s latency. OSS vs Frontier: gpt-4.1 scores +1 index points higher than llama-3.2-3b-instruct, at 0x the per-request cost. Enabling the guardrail layer raises the insurability index by +0 points on average — the premium reduction a safety layer buys.

Threats to validity: LLM-judge scores carry bias (esp. self-preference when a judge grades its own provider — see per-judge columns and Cohen's k); N is modest, so CIs are wide; prompts are English-only and indicative, not a certification.